Fynd AI Intern – Take Home Assessment

Short Report – Submitted by: Vivek Katariya

Task 1 – Predicting Review Ratings Using Prompts

For Task 1, I tried three different prompts to guess the star rating of a Yelp review.
The idea was to see which prompt works better in terms of accuracy and JSON output.

Prompts I Tried

- **v1 Basic Prompt** – very small, straight-forward instruction
- **v2 Guidelines Prompt** – added clear rules for each rating
- **v3 Few-Shot Prompt** – gave a few examples to guide the model

Results

Prompt	Accuracy	JSON Validity
v1_basic	**0.66**	0.75
v2_guidelines	0.47	**0.85**
v3_few_shot	0.53	0.65

Short Discussion

- The **basic prompt** actually worked the best for accuracy because it was simple.
- The **guidelines prompt** produced the cleanest JSON since the rules were strict.

- The **few-shot prompt** was mixed — sometimes good, sometimes not. Overall, simple prompts helped accuracy and detailed prompts helped formatting.

Task 2 – AI Feedback System (Streamlit App)

For Task 2, I built a small Streamlit app with **two dashboards**:

⭐ User Dashboard

- User gives a star rating and review
- AI generates:
  - a polite reply
  - one-line summary
  - 2–3 action suggestions
- Everything gets saved in feedback_data.csv

📊 Admin Dashboard

Admins can see:

- total feedback count
- average rating
- rating distribution chart
- full table of reviews + AI summary + actions
- detailed view for each entry

This makes it easy to understand customer sentiment quickly.

Deployment

The app is deployed on Streamlit Cloud:

👉 https://fynd-ai-assignment-hobvrmv8vtfm3adxngkp8u.streamlit.app/
(Both dashboards are inside the same app through the sidebar.)

Included in Submission

- Task 1 notebook
- Streamlit app code
- requirements.txt
- feedback CSV
- this short report
- live URL